



Network Mesh Implementation with Cilium

Cilium container networking and power of eBPF technology





- Pongsak Khamdee (Tao)
- Senior Systems Engineer
- Nutanix Thailand
- pongsak.khamdee@nutanix.com

- Life: <https://pongsak.online>
- Ex VMware
- Ex Pivotal
- Ex Redhat Consult, Marlo Group
- Ex True, TrueMove, Orange

2



Page Context:

{ type: Book, id: ISBN-10 0-321-83457-7 }

Shop All Departments

Search

Books

Advanced Search

Browse Genres

Bestsellers

New & Future Releases

Paperback


ImageService

Domain-Driven Design and

BookService

books are available for Am

Click to LOOK INSIDE!



Domain-driven Design: Tackling Complexity in the Heart of Software [Hardcover]

Eric Evans (Author)

★★★★☆ (12 customer reviews)

Like (15)

ReviewService

RRP: £46.99

Price: £30.09 & this item is available in the UK with Super Saver Delivery. [See details and](#)

You Save: £16.90 (36%)

PriceService

In stock.

InventoryService

Dispatched from and sold by Amazon.co.uk. Gift-wrap available.

Want guaranteed delivery by Friday, March 23? Order it in the next 21 hours and 9 minutes, and choose Express de

Customers Who Bought This Item Also Bought

OthersAlsoBoughtService

LOOK INSIDE!

Patterns of Enterprise
Application
Architecture... by Martin
Fowler

★★★★☆ (14)

£33.59

ImageService

BookService

Growing Object-
Oriented Software,
Guided by Test... by
S

★★★★☆ (14)

£18.35

ReviewService

PriceService


LOOK INSIDE!

Clean Code: A
Handbook of Agile
Software Craf... by
Robert C. Martin

★★★★☆ (27)

£17.00

LOOK INSIDE!

Enterprise Integration
Patterns: Designing,
Build... by Gregor Hohpe

★★★★☆ (16)

£29.61

LOOK INSIDE!

Continuous Delivery:
Reliable Software
Releases Thr... by Jez
Humble

★★★★☆ (7)

£19.97

LOOK INSIDE!

Refactoring: Improving
the Design of Existing
Co... by Martin Fowler

★★★★☆ (24)

£29.61


Product details

Hardcover: 560 pages


Publisher: Addison Wesley; 1 edition (20 Aug 2003)

BookService


Multi-Cluster deployment



Multi-Cluster deployment




Cluster Mesh




Kubernetes fundamental network requirement

- **Pod-to-Pod Communication without NAT:** Every pod in a Kubernetes cluster receives a unique IP address. The network must allow direct communication between any two pods, regardless of whether they reside on the same node or different nodes, without requiring Network Address Translation (NAT) for internal pod-to-pod traffic.
- **Node-to-Pod Communication:** Agents running on a node, such as system daemons and the Kubelet, must be able to communicate with all pods running on that specific node.
- **Service-to-Pod Communication:** Kubernetes Services provide a stable abstraction layer for accessing pods. The network must enable communication from Services to the underlying pods they expose.
- **External Connectivity (Ingress/Egress):** For applications to interact with the outside world, the network needs to facilitate both ingress (incoming) and egress (outgoing) traffic. This often involves using NAT gateways for pods to access external resources and mechanisms like Ingress controllers or Load Balancers to expose Services to external clients.


Let's get 3 VMs



Create a bridge on each of them




And start the containers




Wait


If all the bridges have same subnet,
Things can get messy



So, let's give them a subnet each



But, how does a packet get routed from etcd1 to etcd2?



There are many ways

So, some smart people got together and
wrote a specification

Which is called
CNI
(Container Networking Interface)

Interface between Container Orchestration Engines and Container Runtimes




Networking

HOST network

10.0.30.0/16



Networking






Networking






Networking






Networking






Networking






Networking






Networking






Networking



HTTP request latency via k8s service (usec)

■ eBPF ■ kube-proxy (iptables)



What is Cilium

- Cilium is an opensource solution, which is based on eBPF technology for cloud-native networking and security and was donated to the CNCF in 2021
- In 2023, Cilium is first on only Graduated CNCF project in cloud-native networking
- Adopted as the default CNI by most managed K8S services and distributions.
- In 2024, Cisco acquired its parent company, Isovalent.
- **Commitment to opensource:** Cisco has stated its commitment to continuing to support and nurture opensource communities for Cilium.




cilium

What is eBPF technology


- **Kernel-level execution:** eBPF compiles programs into a bytecode that runs directly in kernel space, the core of the operating system.
- **Event-driven:** eBPF programs are triggered by specific events, such as system call, network packages, or function call in either the kernel or user space.
- **Dynamic extension:** Allow extension of functionality at runtime without needing to load a kernel module or change the kernel's source code.
- **Key use cases:** Networking (network filtering, lb, networkpolicy), Security (visibility into system calls, real-time threat detection), Observability (collect telemetry, tracing)




Cilium CNI – Encapsulating Routing Mode




Cilium CNI – Native Routing Mode




Cilium Cluster Mesh



Enforce network policy spanning multi cluster



Cilium Cluster Mesh Architecture



Prerequisites

- CIDR ranges in all clusters and all nodes must be non-conflicting and unique IP addresses
- Nodes in all clusters must have IP connectivity between each other using the configured InternalIP for each node. This requirement is typically met by establishing peering or VPN tunnels between the networks of the nodes of each cluster
- The network between clusters must allow inter-cluster communication. The exact ports are documented in the [Firewall Rules](#) section
- Service Load Balancer IP range needs to have a minimum of 2 IP addresses (one for NKP, one for Cluster Mesh)

<https://www.nutanix.dev/2025/09/17/how-to-enable-cilium-cluster-mesh-in-nutanix-kubernetes-platform/>

Cilium Cluster Mesh Configmap

```
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: <NAME>
  namespace: <WORKSPACE_NAMESPACE>
spec:
  topology:
    controlPlane:
      metadata:
        annotations:
          controlplane.cluster.x-k8s.io/skip-kube-proxy: ""
      variables:
        - name: clusterConfig
          value:
            addons:
              cni:
                provider: Cilium
                strategy: HelmAddon
                values:
                  sourceRef:
                    name: cilium-clustermesh-clustera-config
                    kind: ConfigMap
```

```
cat <<EOF | kubectl apply -f -
apiVersion: v1
kind: ConfigMap
metadata:
  name: cilium-clustermesh-clustera-config
  namespace: ${WORKSPACE_NAMESPACE}
data:
  values.yaml: |-
    # Cilium configuration values
    kubeProxyReplacement: true
    k8sServiceHost: "auto"
    tls:
      ca:
        cert: ${CA_CERT}
        key: ${CA_KEY}
    certgen:
      image:
        useDigest: false
    cluster:
      id: 1
      name: cluster-a
    clustermesh:
      apiserver:
        service:
          type: LoadBalancer
          loadBalancerIP: 10.68.252.26 # modify with your own IP
    tls:
      auto:
        enabled: true
        method: cronJob
        schedule: 0 0 1 */4 *
```

```
config:
  clusters:
    - ips:
        - 10.68.252.36 # modify with your own IP
        name: cluster-b
        port: 2379
      enabled: true
      useAPIServer: true
  cni:
    chainingMode: portmap
    exclusive: false
  envoy:
    image:
      useDigest: false
  hubble:
    enabled: true
    relay:
      enabled: true
      image:
        useDigest: false
      priorityClassName: system-cluster-critical
    tls:
      server:
        enabled: true
        mtls: true
  tls:
    auto:
      certValidityDuration: 60
      enabled: true
      method: cronJob
      schedule: 0 0 1 * *
  image:
    useDigest: false
  ipam:
    mode: kubernetes
  operator:
    image:
      useDigest: false
    routingMode: tunnel
    socketLB:
      hostNamespaceOnly: true
    tunnelProtocol: vxlan
```

EOF

Cilium status

```
[nutanix@harbor dev1]$ KUBECONFIG=cluster-a.yaml cilium status --wait
```

```

  Cilium:          OK
  Operator:        OK
  Envoy DaemonSet: OK
  Hubble Relay:    OK
  ClusterMesh:     OK

```

```

DaemonSet      cilium           Desired: 2, Ready: 2/2, Available: 2/2
DaemonSet      cilium-envoy     Desired: 2, Ready: 2/2, Available: 2/2
Deployment      cilium-operator  Desired: 2, Ready: 2/2, Available: 2/2
Deployment      clustermesh-apiserver Desired: 1, Ready: 1/1, Available: 1/1
Deployment      hubble-relay     Desired: 1, Ready: 1/1, Available: 1/1
Containers:
  cilium         Running: 2
  cilium-envoy   Running: 2
  cilium-operator Running: 2
  clustermesh-apiserver Running: 1
  hubble-relay   Running: 1

```

```
Cluster Pods: 18/18 managed by Cilium
```

```
Helm chart version: 1.17.4
```

```
Image versions
cilium           quay.io/cilium/cilium:v1.17.4: 2
cilium-envoy     quay.io/cilium/cilium-envoy:v1.32.6-1746661844-
cilium-operator  quay.io/cilium/operator-generic:v1.17.4: 2
clustermesh-apiserver quay.io/cilium/clustermesh-apiserver:v1.17.4@sl
hubble-relay     quay.io/cilium/hubble-relay:v1.17.4: 1
```

```
[nutanix@harbor dev1]$
```

```
[nutanix@harbor dev1]$ KUBECONFIG=cluster-b.yaml cilium status --wait
```

```

  Cilium:          OK
  Operator:        OK
  Envoy DaemonSet: OK
  Hubble Relay:    OK
  ClusterMesh:     OK

```

```

DaemonSet      cilium           Desired: 2, Ready: 2/2, Available: 2/2
DaemonSet      cilium-envoy     Desired: 2, Ready: 2/2, Available: 2/2
Deployment      cilium-operator  Desired: 2, Ready: 2/2, Available: 2/2
Deployment      clustermesh-apiserver Desired: 1, Ready: 1/1, Available: 1/1
Deployment      hubble-relay     Desired: 1, Ready: 1/1, Available: 1/1
Containers:
  cilium         Running: 2
  cilium-envoy   Running: 2
  cilium-operator Running: 2
  clustermesh-apiserver Running: 1
  hubble-relay   Running: 1

```

```
Cluster Pods: 26/26 managed by Cilium
```

```
Helm chart version: 1.17.4
```

```
Image versions
cilium           quay.io/cilium/cilium:v1.17.4: 2
cilium-envoy     quay.io/cilium/cilium-envoy:v1.32.6-1746661844-0f602c28cb2aa57b29078195049fb257d5b5246c: 2
cilium-operator  quay.io/cilium/operator-generic:v1.17.4: 2
clustermesh-apiserver quay.io/cilium/clustermesh-apiserver:v1.17.4@sha256:0b72f3046cf36ff9b113d53cc61185e893edb5fe728a2c9e561c1083f806453d: 3
hubble-relay     quay.io/cilium/hubble-relay:v1.17.4: 1
```

```
[nutanix@harbor dev1]$
```

Cluster Mesh status

```
[nutanix@harbor dev1]$ KUBECONFIG=cluster-a.yaml cilium clustermesh status --wait
```

```
✓ Service "clustermesh-apiserver" of type "LoadBalancer" found
```

```
✓ Cluster access information is available:
```

```
  - 10.38.19.210:2379
```

```
✓ Deployment clustermesh-apiserver is ready
```

```
i KVStoreMesh is enabled
```

```
✓ All 2 nodes are connected to all clusters [min:1 / avg:1.0 / max:1]
```

```
✓ All 1 KVStoreMesh replicas are connected to all clusters [min:1 / avg:1.0 / max:1]
```

```
🔧 Cluster Connections:
```

```
  - cluster-b: 2/2 configured, 2/2 connected - KVSt
```

```
[nutanix@harbor dev1]$
```

```
[nutanix@harbor dev1]$ KUBECONFIG=cluster-b.yaml cilium clustermesh status --wait
```

```
✓ Service "clustermesh-apiserver" of type "LoadBalancer" found
```

```
✓ Cluster access information is available:
```

```
  - 10.38.19.213:2379
```

```
✓ Deployment clustermesh-apiserver is ready
```

```
i KVStoreMesh is enabled
```

```
✓ All 2 nodes are connected to all clusters [min:1 / avg:1.0 / max:1]
```

```
✓ All 1 KVStoreMesh replicas are connected to all clusters [min:1 / avg:1.0 / max:1]
```

```
🔧 Cluster Connections:
```

```
  - cluster-a: 2/2 configured, 2/2 connected - KVStoreMesh: 1/1 configured, 1/1 connected
```

```
[nutanix@harbor dev1]$
```

Cluster Mesh – Global Service (cluster A)

```
# shared-service.yaml
```

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: shared-service
```

```
---
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deployment
```

```
  namespace: shared-service
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      app: nginx
```

```
  replicas: 2
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: nginx
```

```
    spec:
```

```
      containers:
```

```
      - name: nginx
```

```
        image: nginx:latest
```

```
        ports:
```

```
        - containerPort: 80
```

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
  name: shared-service
```

```
  namespace: shared-service
```

```
  annotations:
```

```
    service.cilium.io/global: "true" # Crucial for Cluster Mesh service d
```

```
spec:
```

```
  selector:
```

```
    app: nginx
```

```
  ports:
```

```
  - protocol: TCP
```

```
    port: 80
```

```
    targetPort: 80
```

```
  type: ClusterIP
```


Cluster Mesh – Global service (cluster B)

```
apiVersion: v1
kind: Service
metadata:
  name: shared-service
  namespace: shared-service
  annotations:
    service.cilium.io/global: "true" # Crucial for Cluster Mesh service d
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
  type: ClusterIP
```

Cluster Mesh – Test Pod (cluster B)

```
# client-pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: client-pod
  namespace: shared-service
spec:
  containers:
  - name: client
    image: curlimages/curl:latest
    command: ["tail", "-f", "/dev/null"] # Keep the pod running
```

Cluster Mesh – Test Pod (cluster B)

```
[nutanix@harbor dev1]$ kubectl -n shared-service exec client-pod -- curl shared-service.shared-service
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100   615  100   615    0     0  244338      0 --:--:-- --:--:-- --:--:-- 307500
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>


<p><em>Thank you for using nginx.</em></p>
</body>
</html>
[nutanix@harbor dev1]$
```



One More Thing...


Service Mesh

- A dedicated infrastructure layer for managing service-to-service communication.
- Data Plane: The layer responsible for handling actual application traffic (Proxies that handle traffic)
- Control plane: It configures and manages the data plane. Management brain that tells proxies what to do.




How it works (Sidecar)

- Each pod gets a proxy container e.g. Envoy.
- All inbound/outbound traffic goes through the proxy
- Service Mesh control plane **istiod** configures proxies
- Pilot + Citadel + Galley merged into istiod.



How it works (Ambient Mode)

- No sidecars no maintenance overhead
- Traffic intercepted by proxy runs once per node DaemonSet, L4 secure overlay
- ztunnel per-node proxy that transparently intercepts pod traffic
- It tunnels traffic using HBONE (HTTP Based Overlay Network). Multiple TCP streams multiplex over a single mTLS connection.



When to use

- Many microservices needing secure, reliable communication
- Want observability without adding code
- Need fine-grained traffic control
- expect service count to grow fast
- run multi-cluster / multi-cloud
- want security, routing, and monitoring without touching app code
- want easier certificate rotation


When not to use

- Have a small number of services, few services (say <20–30)
- Communication is mostly internal & trusted, stays inside one trusted cluster
- Don't urgently need centralized traffic policies or observability
- Performance overhead is a major concern
- Fine maintaining mTLS in code & cert rotation manually


One More Thing...

Gateway API



Service Exposure




ClusterIP



LoadBalancer



NodePort



Ingress

Ingress is lagacy


```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: basic-ingress
  namespace: default
spec:
  ingressClassName: cilium
  rules:
  - http:
      paths:
      - backend:
          service:
            name: details
            port:
              number: 9080
        path: /details
        pathType: Prefix
      - backend:
          service:
            name: productpage
            port:
              number: 9080
        path: /
        pathType: Prefix
```

Gateway API

```
apiVersion: gateway.networking.k8s.io/v1
kind: Gateway
metadata:
  name: my-gateway
spec:
  gatewayClassName: cilium
  listeners:
  - protocol: HTTP
    port: 80
    name: web-gw
  allowedRoutes:
    namespaces:
      from: Same
```

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: http-app-1
spec:
  parentRefs:
  - name: my-gateway
    namespace: default
  rules:
  - matches:
    - path:
        type: PathPrefix
        value: /details
    backendRefs:
    - name: details
      port: 9080
  - matches:
    - headers:
      - type: Exact
        name: magic
        value: foo
      queryParams:
      - type: Exact
        name: great
        value: example
    path:
      type: PathPrefix
      value: /
    method: GET
  backendRefs:
  - name: productpage
    port: 9080
```

Gateway API Use Cases



Gateway API Use Cases

NUTANIX

“Network Mesh Implementation with Cilium”

ร่วมทำแบบสอบถามเพื่อรับของที่ระลึกจากนุทานิคส์



กระเป๋อเนกประสงค์ / เป๋าสายไฟ

NUTANIX

Thank you